

ECE 219 Project 2: Clustering

Chuang Yu (305629107), Hongyi Chen (705186099), Zichun Chai (505625207)

Question 1:

In this project, we explored the concept of clustering. Clustering algorithms are unsupervised methods for finding groups of data points that have similar representations in a feature space and it differs from classification that we did in the last project since it does not have any a priori labeling of data points.

In this project, we found proper representation of data, performed K-means, HDBscan, and DBscan on the dataset and evaluated the result of clustering. Lastly, we tried different preprocessing techniques that might increase the performance of clustering.

We worked with “20 Newsgroups” dataset that we already explored in Project 1. It is a collection of approximately 20,000 documents, partitioned (nearly) evenly across 20 different newsgroups, each corresponding to a different category (topic). Each topic can be viewed as a “class”.

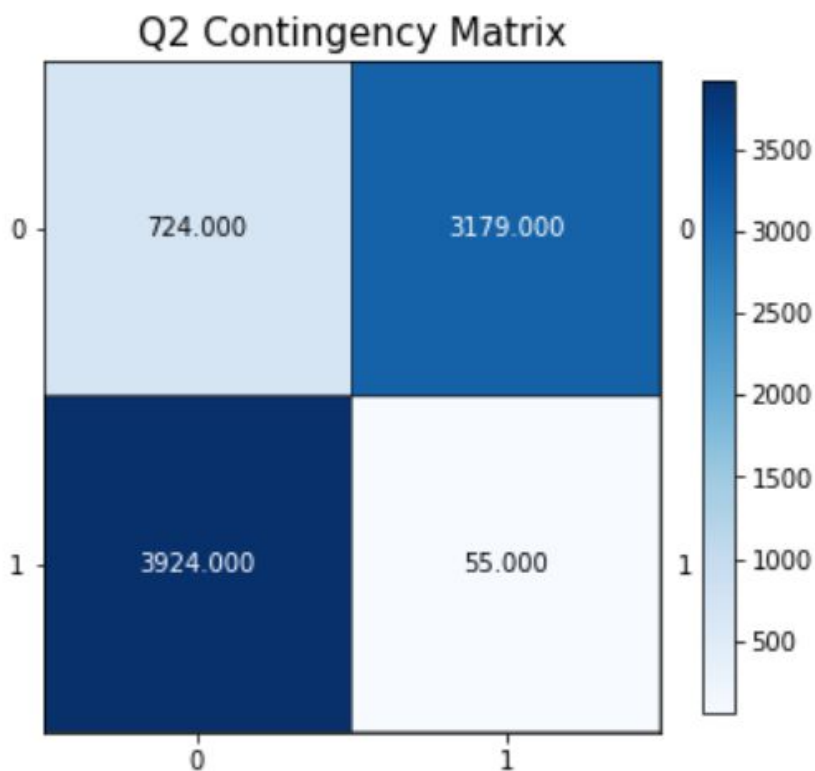
We followed the steps in Project 1, and transformed the documents into TF-IDF vectors. This measure takes into account count of words in the document, as normalized by a certain function of the frequency of the individual words in the whole corpus. Also, we used `min_df = 3` and `stop_words='english'` in `CountVectorizer` without any stemming or lemmatization, and removed headers and footers in 20newsgroups. As for the 20newsgroups, we specified 8 categories and `random_state` to be 0. As a result, the dimensions of the TF-IDF matrix is:

```
<class 'scipy.sparse.csr.csr_matrix'>  
(7882, 23522)
```

Question 2:

In this problem, we applied K-means clustering with $k = 2$ using the TF-IDF data from Q1. As for the parameters in the K-means class in sklearn, we specified that `random_state=0`, `max_iter=1000` and `n_init=30`. The `random_state` parameter denoted the random number generation for centroid initialization. We needed to use an integer to make the randomness deterministic. The `max_iter` parameter represented the maximum number of iterations of the k-means algorithm for a single run. And the `n_init` parameter denoted the number of times that the k-means algorithm will be run with different centroid seeds. These parameters can improve the performance at a cost of computation power.

Given the clustering result and ground truth labels, contingency table A is the matrix whose entries A_{ij} is the number of data points that belong to both the class C_i the cluster K_j . While plotting the contingency table, we utilized the provided `plotmat.py` file to visualize the matrix.



Question 3:

In order to evaluate clustering results, there are various measures for a given partition of the data points with respect to the ground truth. We adopted the measures homogeneity score, completeness score, V-measure, adjusted Rand score and adjusted mutual info score, all of which could be calculated by the corresponding functions provided in sklearn.metrics. Homogeneity is a measure of how “pure” the clusters are. On the other hand, a clustering result satisfies completeness if all data points of a class are assigned to the same cluster. And the V-measure is then defined to be the harmonic average of homogeneity score and completeness score. The adjusted Rand Index is similar to accuracy measure, which computes similarity between the clustering labels and ground truth labels. And the adjusted mutual information score measures the mutual information between the cluster label distribution and the ground truth label distributions. We have utilized these 5 measures for the K-means clustering results and the report is shown below:

```
Homogeneity Score: 0.5809557696057134
Completeness Score: 0.5948001867166814
V-measure Score: 0.5877964697988065
Adjusted Rand Index: 0.6436957246204892
Adjusted Mutual Information Score: 0.587758287733538
```

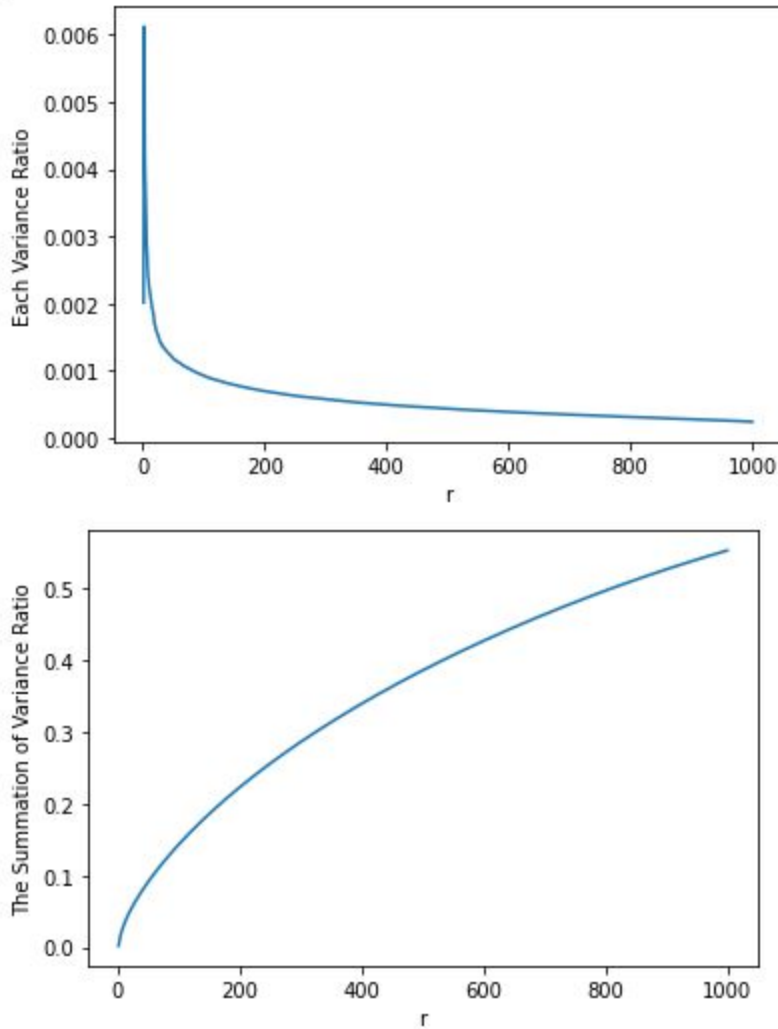
All the above 5 measures range between 0 and 1 and we wanted them to be as close to 1 as possible. Therefore, we would try different methods for dimensionality reduction, data transformation, and clustering later in this project.

Question 4:

From the last problem we can see that high dimensional sparse TF-IDF vectors did not yield a good clustering result. One of the reasons is that in a high-dimensional space, the Euclidean distance is not a good metric anymore, in the sense that the distances between data points tends to be almost the same. Furthermore, K-means clustering is limited since it is good at round-shaped clusters without unequal variances. Therefore, in this problem, we tried to reduce the dimension of our data before clustering and applying other transformations.

Firstly, we investigated on what ratio of the variance of the original data is retained after the dimensionality reduction. Reconstructing the matrix with the truncated SVD representation, we plotted the percent of variance the top r principal components could retain vs r for $r = 1$ to 1000 by using `explained_variance_ratio_` of TruncatedSVD objects.

The plots for both each variance ratio and the sum of variance ratio are shown below:

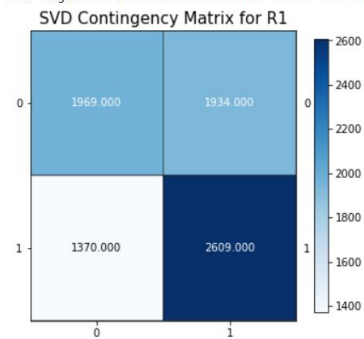


Question 5:

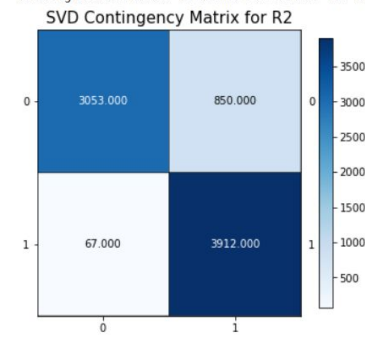
We used two methods to reduce the dimension of the data which were Singular Value Decomposition (SVD) and Non-negative Matrix Factorization (NMF). We swept over the dimension parameters $r = [1, 2, 3, 5, 10, 20, 50, 100, 300]$ of each method, and based on the resulting homogeneity score, completeness score, V-measure, adjusted Rand score, adjusted mutual info score, and contingency matrix, we chose the the r that yielded the best results in terms of clustering purity metrics.

The results for SVD are shown below:

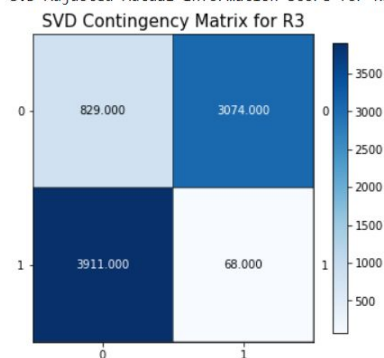
SVD Homogeneity Score for R1 0.019034591198472048
SVD Completeness Score for R1 0.019360460836506627
SVD V-measure Score for R1 0.01919614314371571
SVD Adjusted Rand Index for R1 0.02600277723478503
SVD Adjusted Mutual Information Score for R1 0.01910558756050995



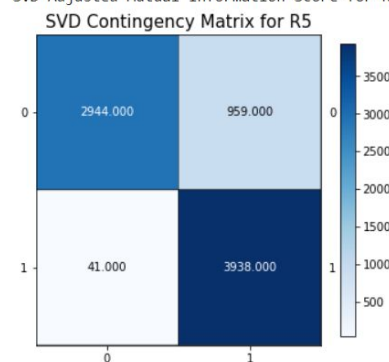
SVD Homogeneity Score for R2 0.5318527560716836
SVD Completeness Score for R2 0.5491344627077224
SVD V-measure Score for R2 0.5403554683557176
SVD Adjusted Rand Index for R2 0.588725244819119
SVD Adjusted Mutual Information Score for R2 0.5403127144748595



SVD Homogeneity Score for R3 0.5377651779879102
SVD Completeness Score for R3 0.5542779518154466
SVD V-measure Score for R3 0.5458967201533174
SVD Adjusted Rand Index for R3 0.5965400007596544
SVD Adjusted Mutual Information Score for R3 0.5458545177246293

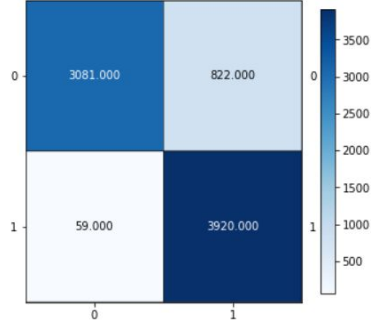


SVD Homogeneity Score for R5 0.517036422502093
SVD Completeness Score for R5 0.5401602853893641
SVD V-measure Score for R5 0.5283454620142469
SVD Adjusted Rand Index for R5 0.5568446089604012
SVD Adjusted Mutual Information Score for R5 0.5283013366507997



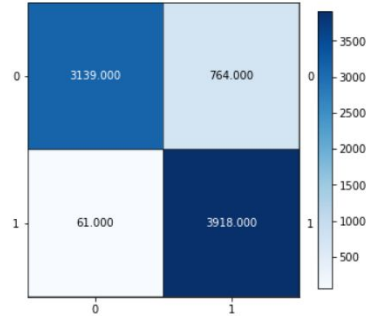
0 1
SVD Homogeneity Score for R10 0.5460921220112308
SVD Completeness Score for R10 0.5629480411958336
SVD V-measure Score for R10 0.5543919879505773
SVD Adjusted Rand Index for R10 0.6028289245085969
SVD Adjusted Mutual Information Score for R10 0.5543505718649027

SVD Contingency Matrix for R10



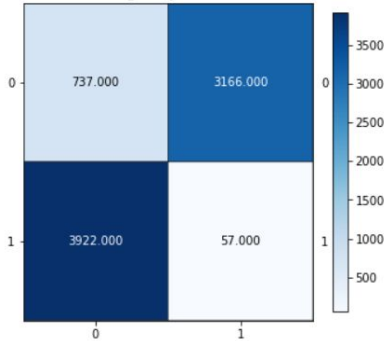
0 1
SVD Homogeneity Score for R20 0.563426312706913
SVD Completeness Score for R20 0.5782222768022195
SVD V-measure Score for R20 0.5707284156217397
SVD Adjusted Rand Index for R20 0.6250997063911322
SVD Adjusted Mutual Information Score for R20 0.5706886059370165

SVD Contingency Matrix for R20



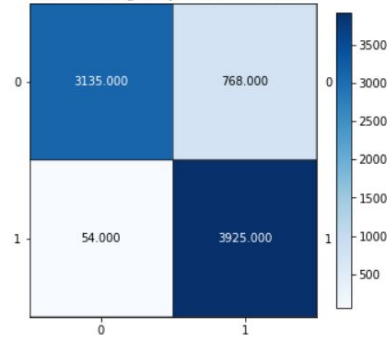
0 1
SVD Homogeneity Score for R50 0.57516931569769
SVD Completeness Score for R50 0.5893200079451415
SVD V-measure Score for R50 0.5821586833212198
SVD Adjusted Rand Index for R50 0.6376018619358004
SVD Adjusted Mutual Information Score for R50 0.5821199646013246

SVD Contingency Matrix for R50



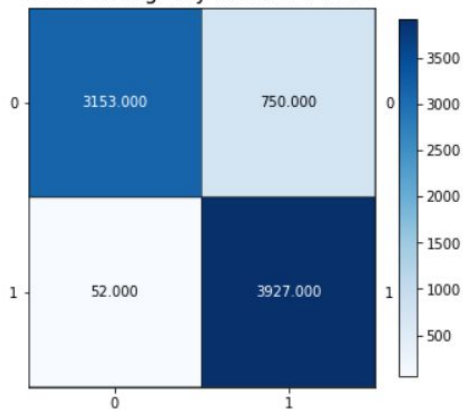
0 1
SVD Homogeneity Score for R100 0.5670330934635874
SVD Completeness Score for R100 0.5823852599054565
SVD V-measure Score for R100 0.5746066513447411
SVD Adjusted Rand Index for R100 0.6263041993292062
SVD Adjusted Mutual Information Score for R100 0.5745671858723812

SVD Contingency Matrix for R100



0 1
SVD Homogeneity Score for R300 0.5744113629390132
SVD Completeness Score for R300 0.5892858802537997
SVD V-measure Score for R300 0.5817535576669169
SVD Adjusted Rand Index for R300 0.6343636646220453
SVD Adjusted Mutual Information Score for R300 0.5817147772526922

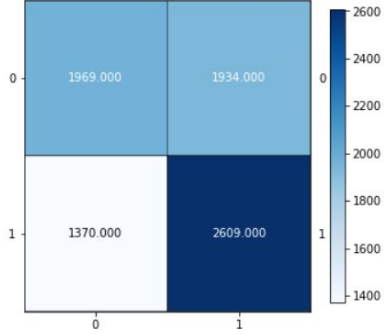
SVD Contingency Matrix for R300



The results for NMF are shown below:

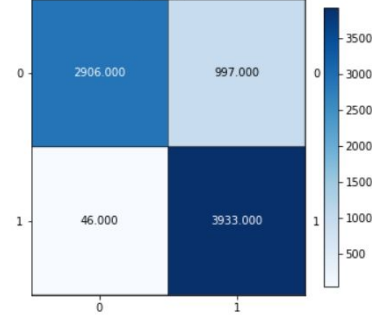
NMF Homogeneity Score for R1 0.019034591198472048
 NMF Completeness Score for R1 0.019360460836506627
 NMF V-measure Score for R1 0.01919614314371571
 NMF Adjusted Rand Index for R1 0.02600277723478503
 NMF Adjusted Mutual Information Score for R1 0.01910558756050995

NMF Contingency Matrix for R1



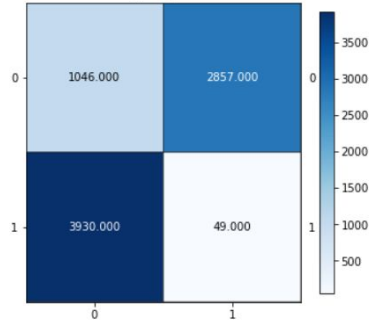
NMF Homogeneity Score for R2 0.5022478335120428
 NMF Completeness Score for R2 0.5263843051392446
 NMF V-measure Score for R2 0.5140328926482289
 NMF Adjusted Rand Index for R2 0.5406769711846864
 NMF Adjusted Mutual Information Score for R2 0.51398735739796

NMF Contingency Matrix for R2



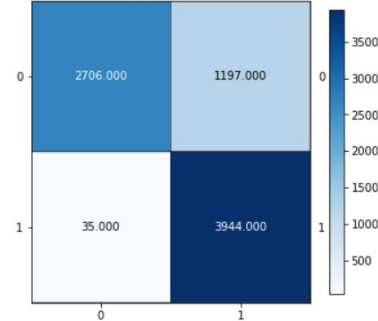
NMF Homogeneity Score for R3 0.48609683035120665
 NMF Completeness Score for R3 0.5118300127229015
 NMF V-measure Score for R3 0.49863163535474486
 NMF Adjusted Rand Index for R3 0.521443534760384
 NMF Adjusted Mutual Information Score for R3 0.4985845503268862

NMF Contingency Matrix for R3



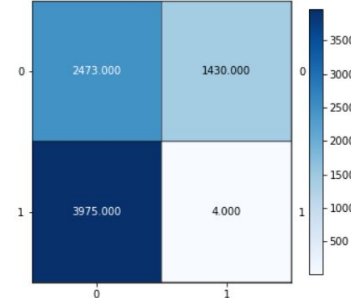
NMF Homogeneity Score for R5 0.45500660137619775
 NMF Completeness Score for R5 0.4881474012058305
 NMF V-measure Score for R5 0.47099474610769254
 NMF Adjusted Rand Index for R5 0.4724383437991488
 NMF Adjusted Mutual Information Score for R5 0.47094461230034934

NMF Contingency Matrix for R5



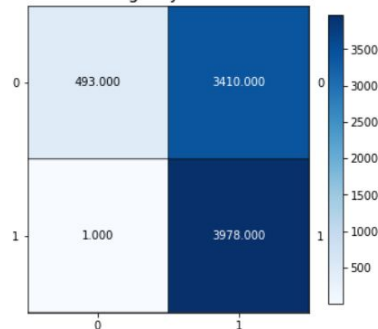
NMF Homogeneity Score for R10 0.20915909036595745
 NMF Completeness Score for R10 0.30563865746695934
 NMF V-measure Score for R10 0.248358132278426
 NMF Adjusted Rand Index for R10 0.1378993331182748
 NMF Adjusted Mutual Information Score for R10 0.24827640731705972

NMF Contingency Matrix for R10



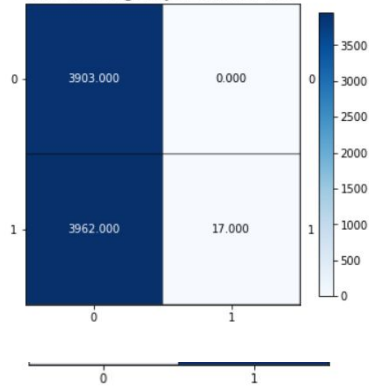
NMF Homogeneity Score for R20 0.06529404662847932
 NMF Completeness Score for R20 0.1931811334155371
 NMF V-measure Score for R20 0.09759991602155817
 NMF Adjusted Rand Index for R20 0.0179867240535735
 NMF Adjusted Mutual Information Score for R20 0.09747631727528892

NMF Contingency Matrix for R20



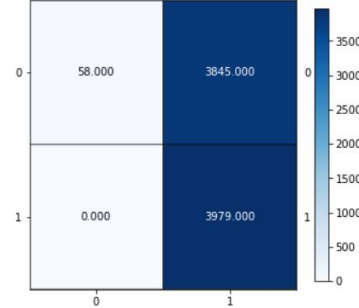
0 1
 NMF Homogeneity Score for R50 0.0021303952018709594
 NMF Completeness Score for R50 0.09591019434272455
 NMF V-measure Score for R50 0.004168204593370178
 NMF Adjusted Rand Index for R50 -6.487642756916324e-05
 NMF Adjusted Mutual Information Score for R50 0.0039841010495616416

NMF Contingency Matrix for R50



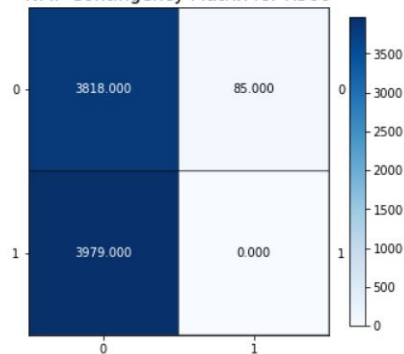
0 1
 NMF Homogeneity Score for R100 0.007502018996706376
 NMF Completeness Score for R100 0.11959893945551682
 NMF V-measure Score for R100 0.014118438235357446
 NMF Adjusted Rand Index for R100 0.0004994593822397486
 NMF Adjusted Mutual Information Score for R100 0.01394709643405834

NMF Contingency Matrix for R100



0 1
 NMF Homogeneity Score for R300 0.011022010519163693
 NMF Completeness Score for R300 0.12823289117096448
 NMF V-measure Score for R300 0.02029923913966382
 NMF Adjusted Rand Index for R300 0.000879750035061344
 NMF Adjusted Mutual Information Score for R300 0.02013308476943905

NMF Contingency Matrix for R300



From the results shown above, we could tell that $r = 50$ is good for SVD, and $r = 2$ is good for NMF, and we also observed that there is a trade-off between the information preservation, and better performance of K-means in lower dimensions. So we would use these two values in the following questions as the best case.

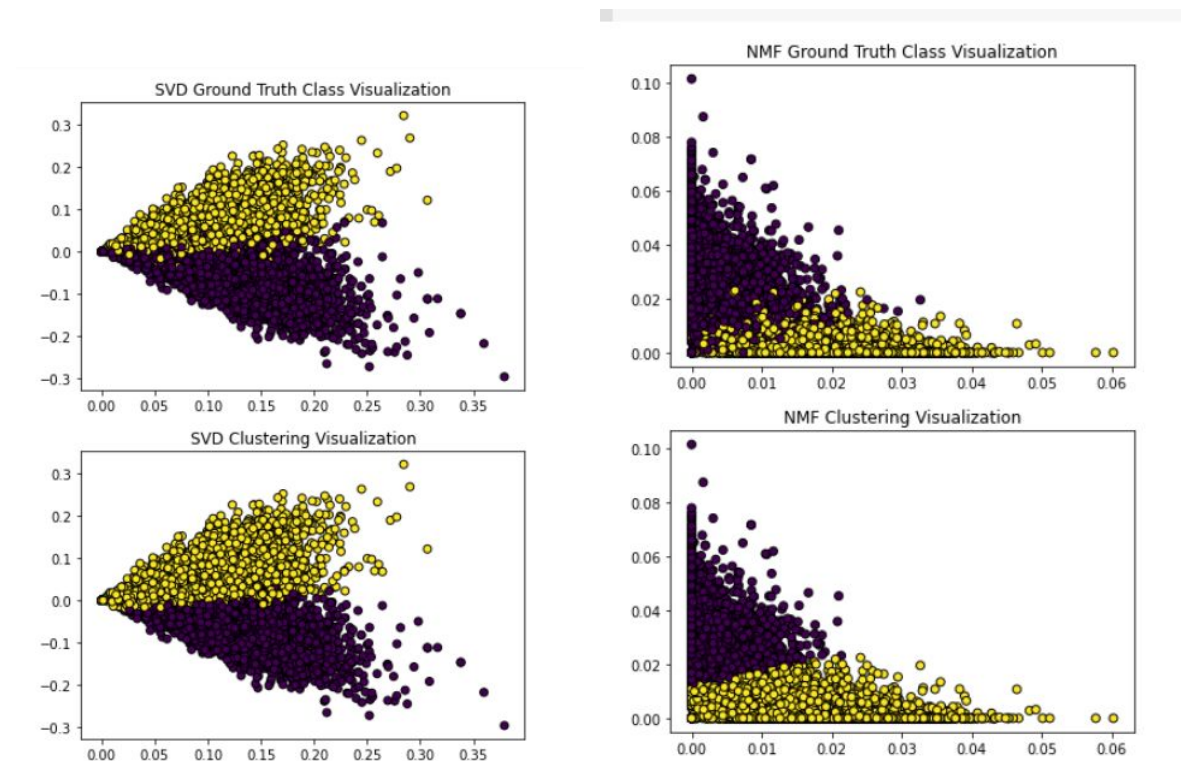
Question 6:

We did observe this non-monotonic behavior of the measures as we increased r . There are mainly two reasons for this. Firstly, if r is too small, we do not have enough information for the TF-IDF vectors and the later clustering process. Secondly, if r is too big, high dimensional sparse TF-IDF vectors do not yield a good clustering result because the Euclidean distance tends to be

almost the same in a high-dimensional space. These two effects would greatly harm the performance of K-means clustering and lead to the non-monotonic behavior.

Question 7:

As for visualization, we could visualize the results by projecting the dim-reduced data points onto a 2-D plane with SVD, and coloring the points according to ground truth class label and clustering label respectively. We chose $r = 50$ for SVD and $r = 2$ for NMF from the results that we have derived from Q5. The plots are shown below:



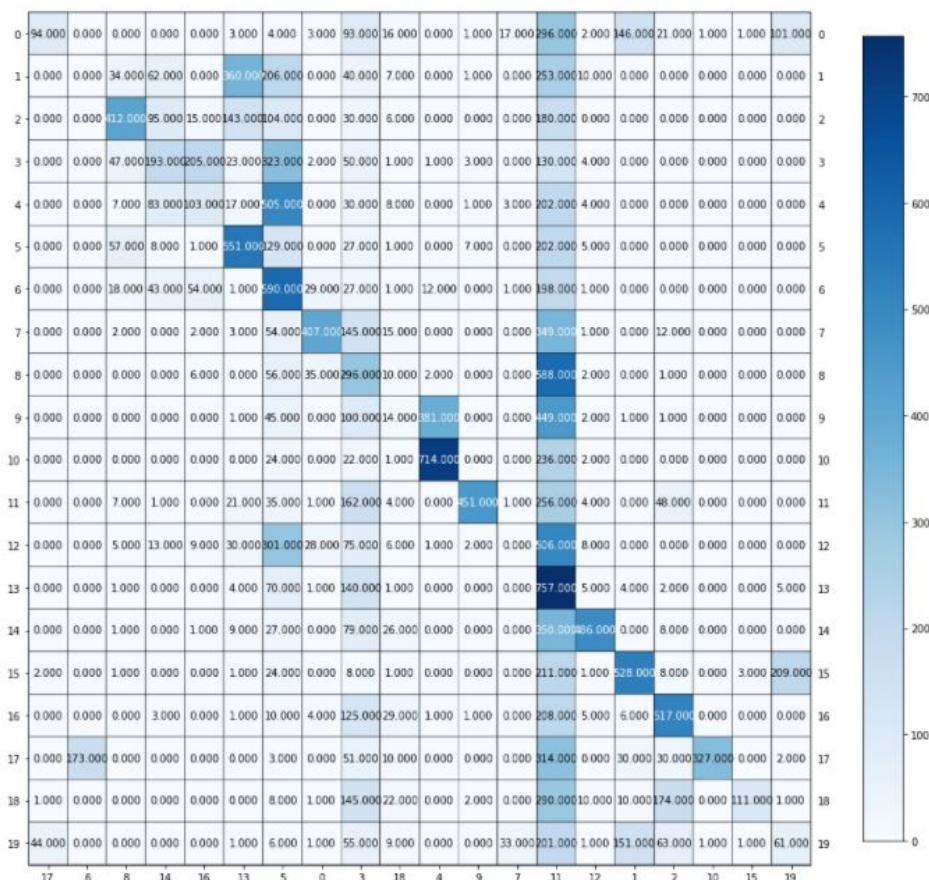
Question 8:

Generally speaking, we observed a boundary between two classes and they were separated clearly. And the ground truth class label plots share a lot in common with the clustering label plots in both dimensionality reduction techniques. However, we still observed some outliers, and therefore we would still need to improve our techniques in dimensionality reduction and clustering for better performance.

Question 9:

In the previous questions, we have been dealing with a relatively simple clustering task with only two well-separated classes “comp” and “rec” for only 8 categories. In this problem, we would cluster for the entire 20 categories in the 20newsgroups dataset. We loaded the documents with the same configuration as in Q1 but for all the 20 categories. We utilized CountVectorizer and TF-IDF for feature extraction, and SVD ($r = 50$) for dimensionality reduction, and at last performed K-means clustering with $n_components = 20$. Also, there is a mismatch between cluster labels and class labels, and as a result, the high-value entries of the $20 * 20$ contingency matrix can be scattered around, making it messy to inspect, even if the clustering result is not bad. Therefore, we used `scipy.optimize.linear_sum_assignment` to identify the best-matching cluster-class pairs, and permute the columns of the contingency matrix accordingly. The contingency matrix and the five clustering metrics are shown below:

```
20 Categories all_dataset SVD Homogeneity Score for R50 0.3458766189788096
20 Categories all_dataset SVD Completeness Score for R50 0.43446715637092265
20 Categories all_dataset SVD V-measure Score for R50 0.38514315318415665
20 Categories all_dataset SVD Adjusted Rand Index for R50 0.11042604980495854
20 Categories all_dataset SVD Adjusted Mutual Information Score for R50 0.3829063568650045
20 Categories all_dataset best SVD with R = 50
```

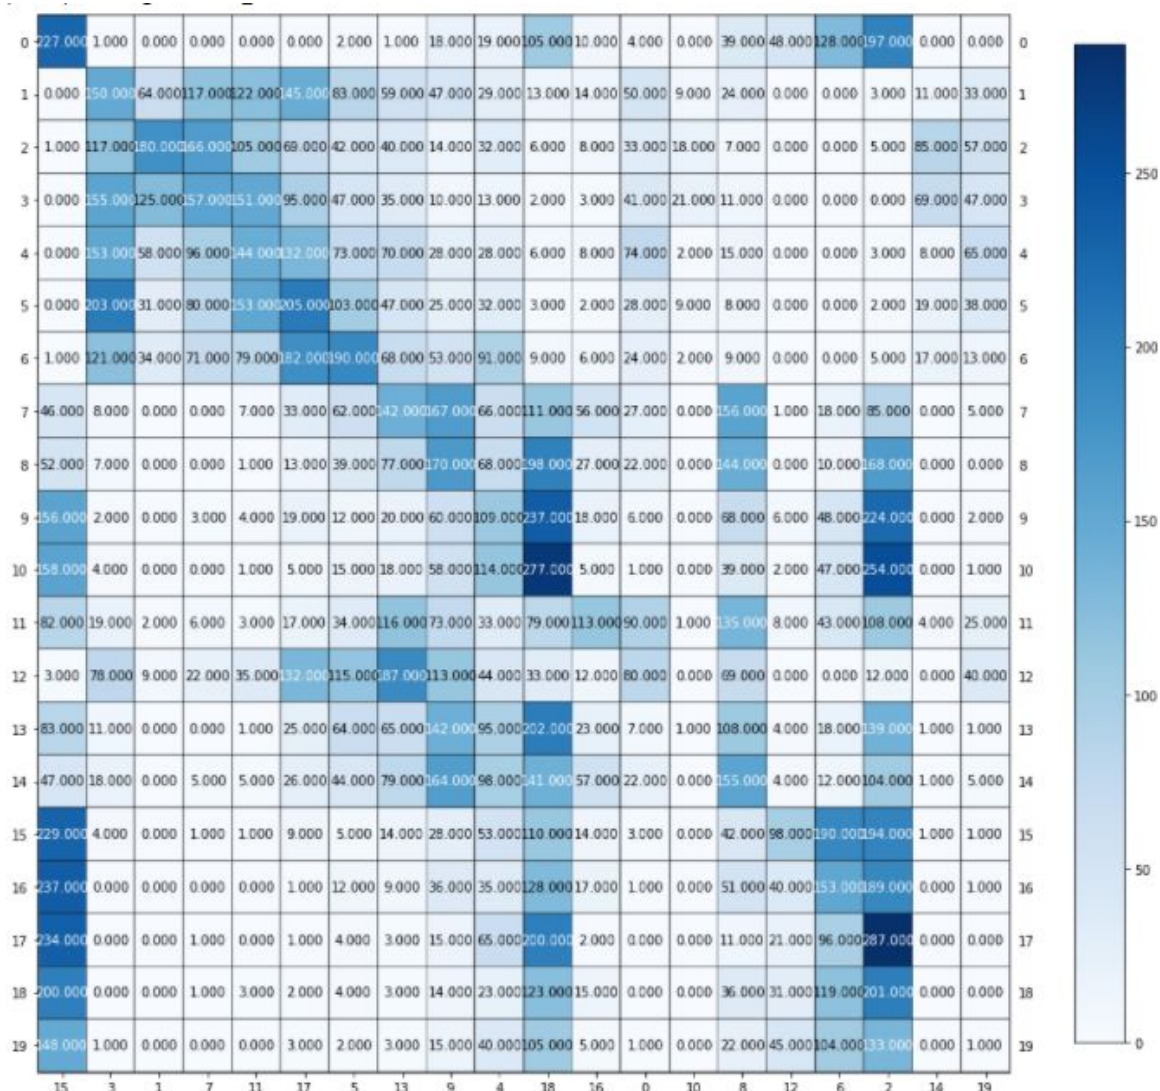


Question 10:

NMF uses Frobenius norm as its cost function by default. And we decided to try another choice which was Kullback-Leibler divergence by specifying `beta_loss = "kullback-leibler"`, `solver = "mu"`, and observed whether it helped with the clustering of our text data. The results for both case with and without KL are shown below:

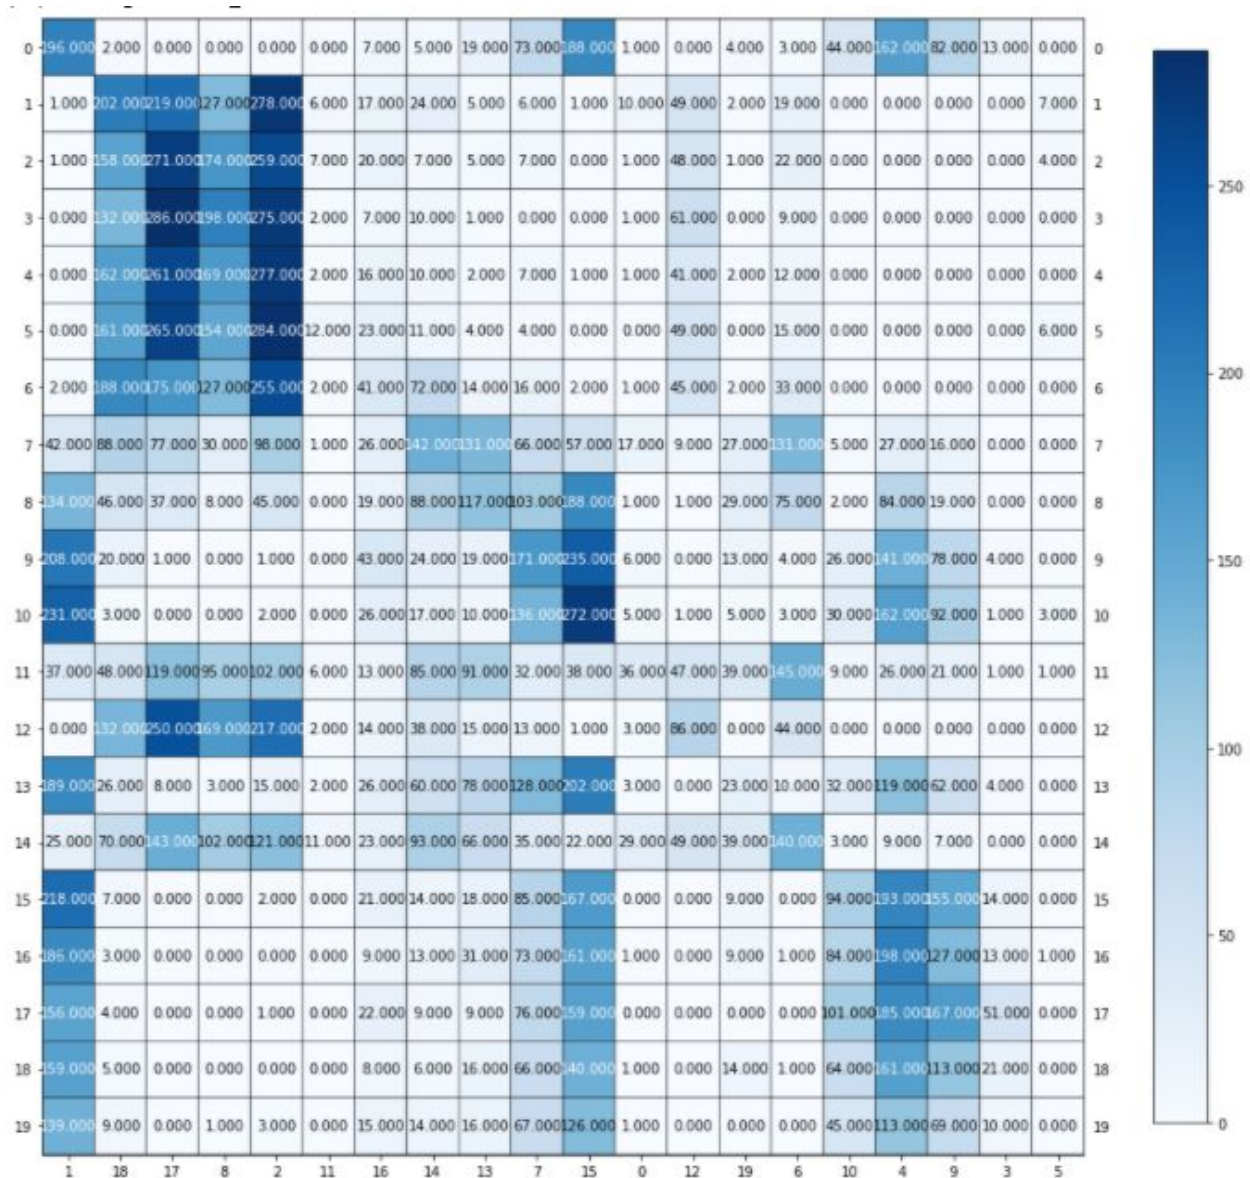
Without KL:

```
(NO KL)20 Categories all_dataset NMF Homogeneity Score for R2 0.19097482490598844
(NO KL)20 Categories all_dataset NMF Completeness Score for R2 0.20452326538392374
(NO KL)20 Categories all_dataset NMF V-measure Score for R2 0.19751698303910672
(NO KL)20 Categories all_dataset NMF Adjusted Rand Index for R2 0.057576439812421415
(NO KL)20 Categories all_dataset NMF Adjusted Mutual Information Score for R2 0.19482165213369498
(NO KL)20 Categories all_dataset best NMF with R = 2
```



With KL:

```
(KL)20 Categoies all_dataset NMF Homogeneity Score for R2 0.20217506002873892
(KL)20 Categoies all_dataset NMF Completeness Score for R2 0.2260432455956135
(KL)20 Categoies all_dataset NMF V-measure Score for R2 0.21344396606656976
(KL)20 Categoies all_dataset NMF Adjusted Rand Index for R2 0.06566950852631316
(KL)20 Categoies all_dataset NMF Adjusted Mutual Information Score for R2 0.21071407416901536
(KL)20 Categoies all_dataset best NMF with R = 2
```



From the results we could see that the case with Kullback-Leibler divergence has a better performance than the case without KL divergence in all five metrics. Therefore, it is clear that

Kullback-Leibler divergence is a better cost function than Frobenius norm because it improves the performance.

Question 11:

In this question we used UMAP to reduce the dimensionality of the 20 categories of the TF-IDF matrix with two settings of metric = 'euclidean' and metric = 'cosine'. The plots below show the results for metric = 'euclidean'. Best n_component is 3.

homogeneity_score 0.008888883402648982

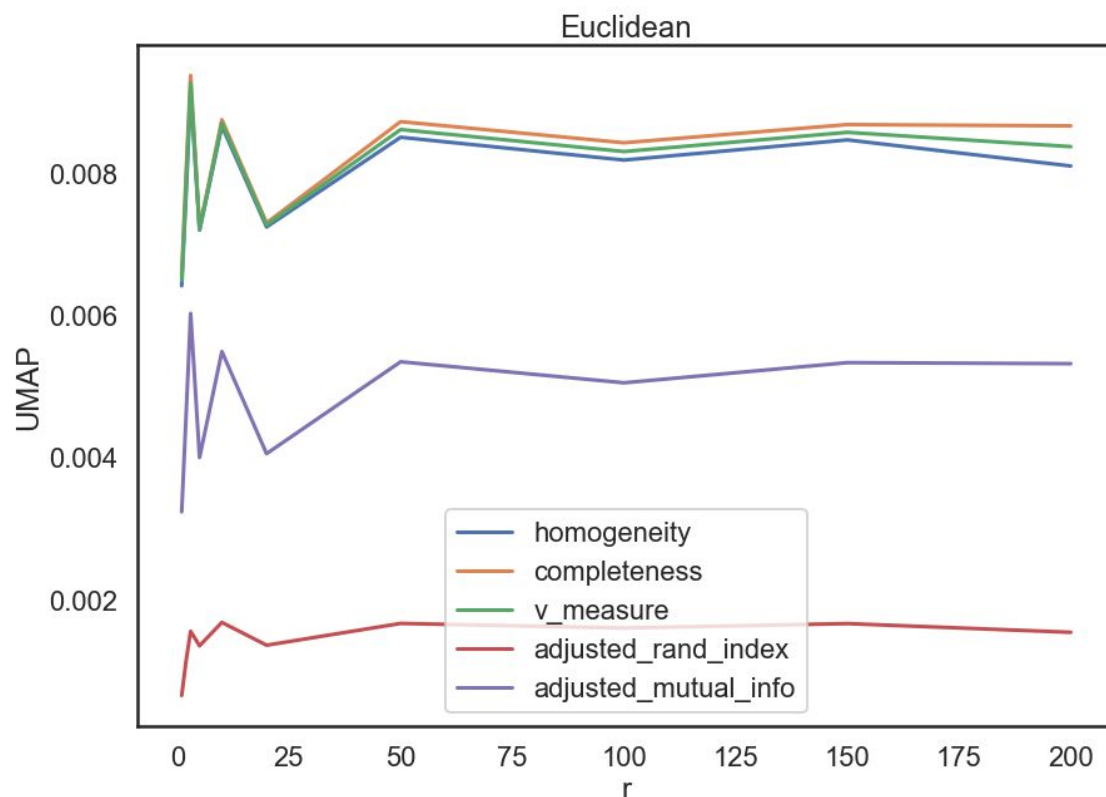
completeness_score 0.009163902973211941

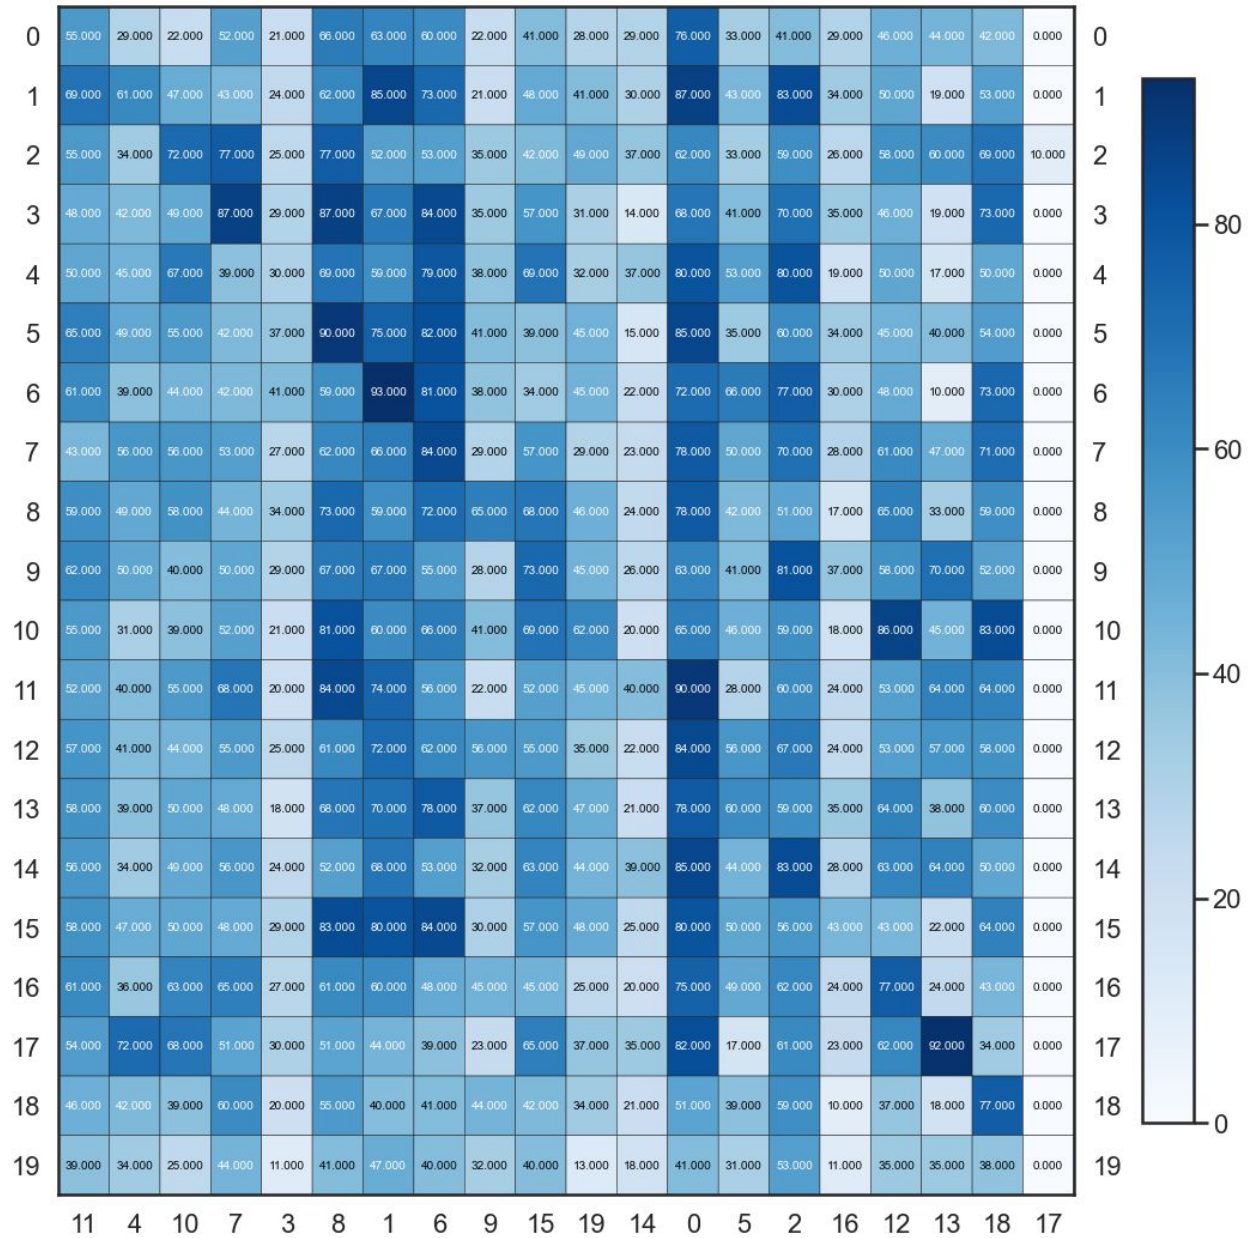
v_measure_score 0.00902429833778883

adjusted_rand_score 0.0014956927303962568

adjusted_mutual_info_score 0.005773132022734455

Contingency matrix with Euclidean





cos

Best n_component is 50,

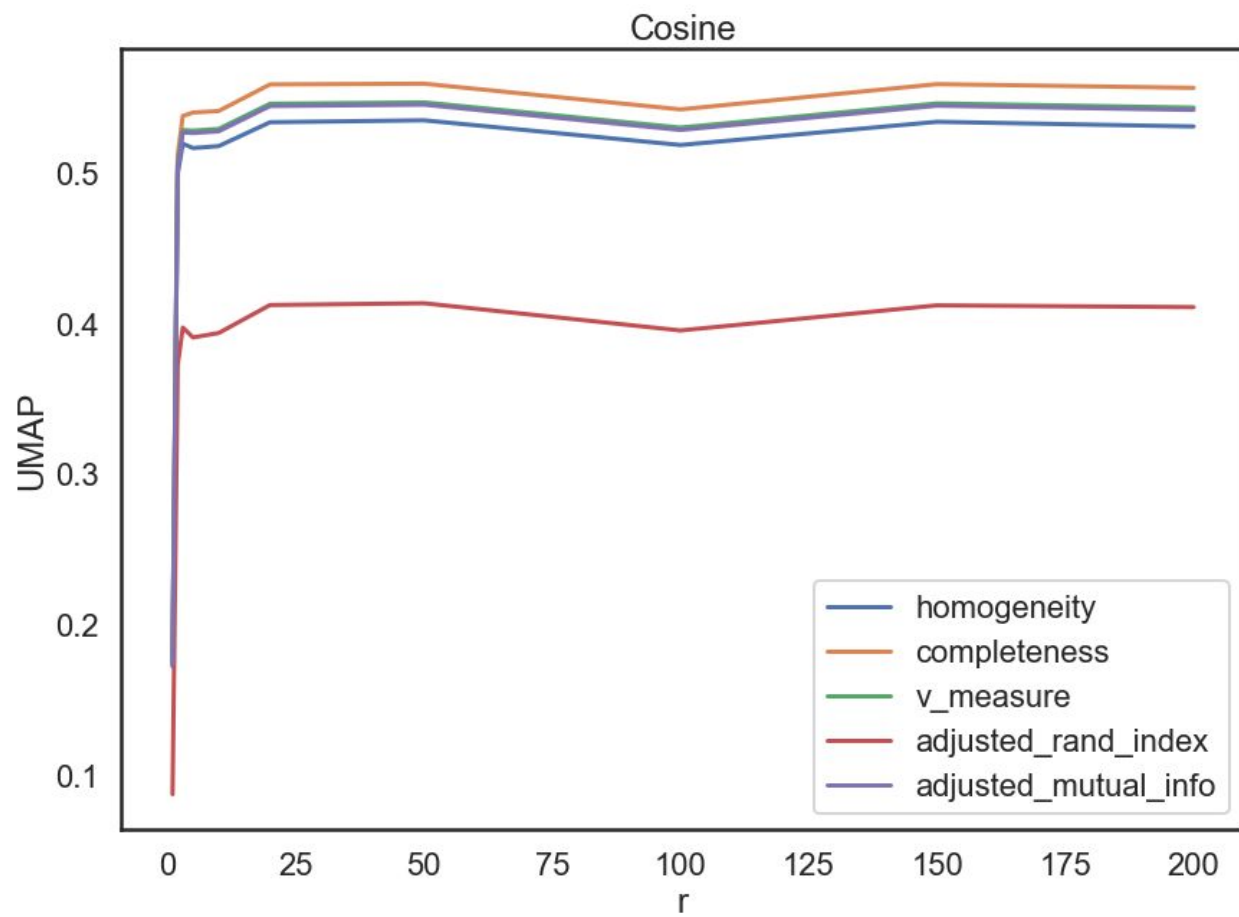
Homogeneity Score 0.5360492927573465

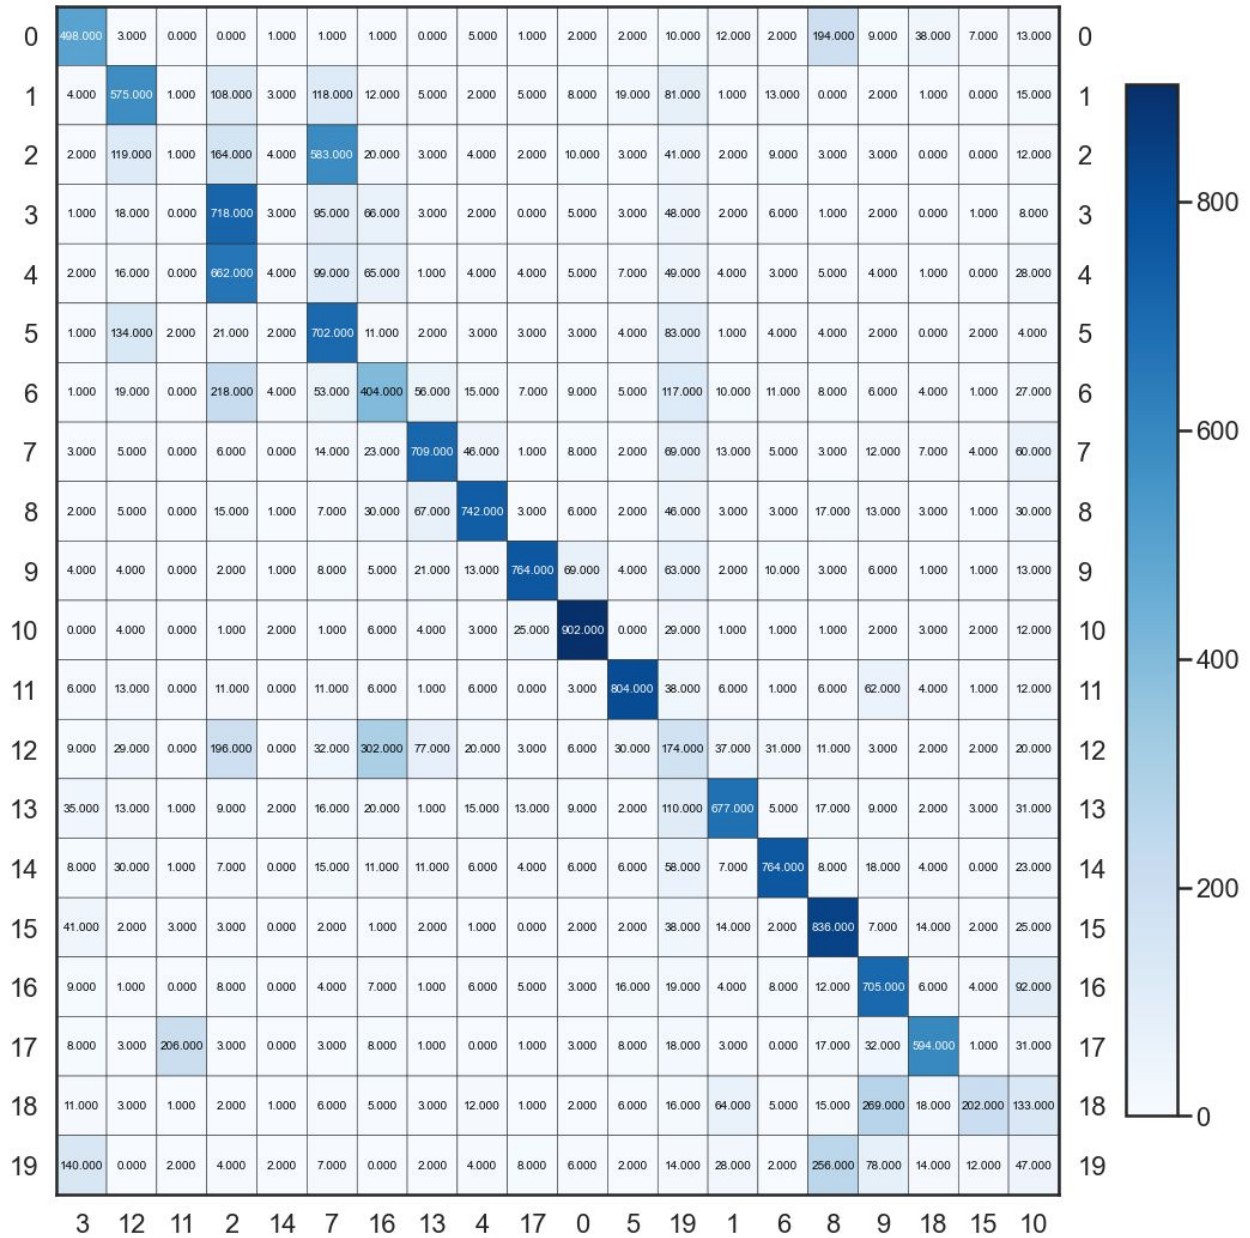
Completeness Score 0.5622405377922715

V-measure Score 0.5488326200602862

Adjusted Rand Index 0.41280179805176986

Adjusted Mutual Information Score 0.5473274512508833





Question 12:

From the contingency matrices above, we can see that the Euclidean matrix is more prone to be confused while from the Cosine confusion matrix it is cleaner and less probable to be confused. The result matched our expectation since the Euclidean metric belongs to Minkowski style metrics and the Cosine metric belongs to angular and correction metrics. Those two set of metrics under UMAP instructions do have the behaviors we saw in our test cases.

Question 13:

For this question, we get following outputs:

“Ward”

Homogeneity Score 0.5017633893942085

Completeness Score 0.5266759659173146

V-measure Score 0.5139179406278006

Adjusted Rand Index 0.3760639944178731

Adjusted Mutual Information Score 0.5122957758731398

“single”

Homogeneity Score 0.020983070656432658

Completeness Score 0.3605274423617933

V-measure Score 0.0396580043722101

Adjusted Rand Index 0.0004612547204987331

Adjusted Mutual Information Score 0.034176765481306826

We can see that the linkage single does not work well, all inputs are clustered wrongly to zero.

Question 14:

For Dbscan, we tested metrics, min_samples and eps these three parameters. The best we get is with eps = 0.55, min_samples = 100 and metric is euclidean

min_sample: 100

Homogeneity 0.4666210225402587

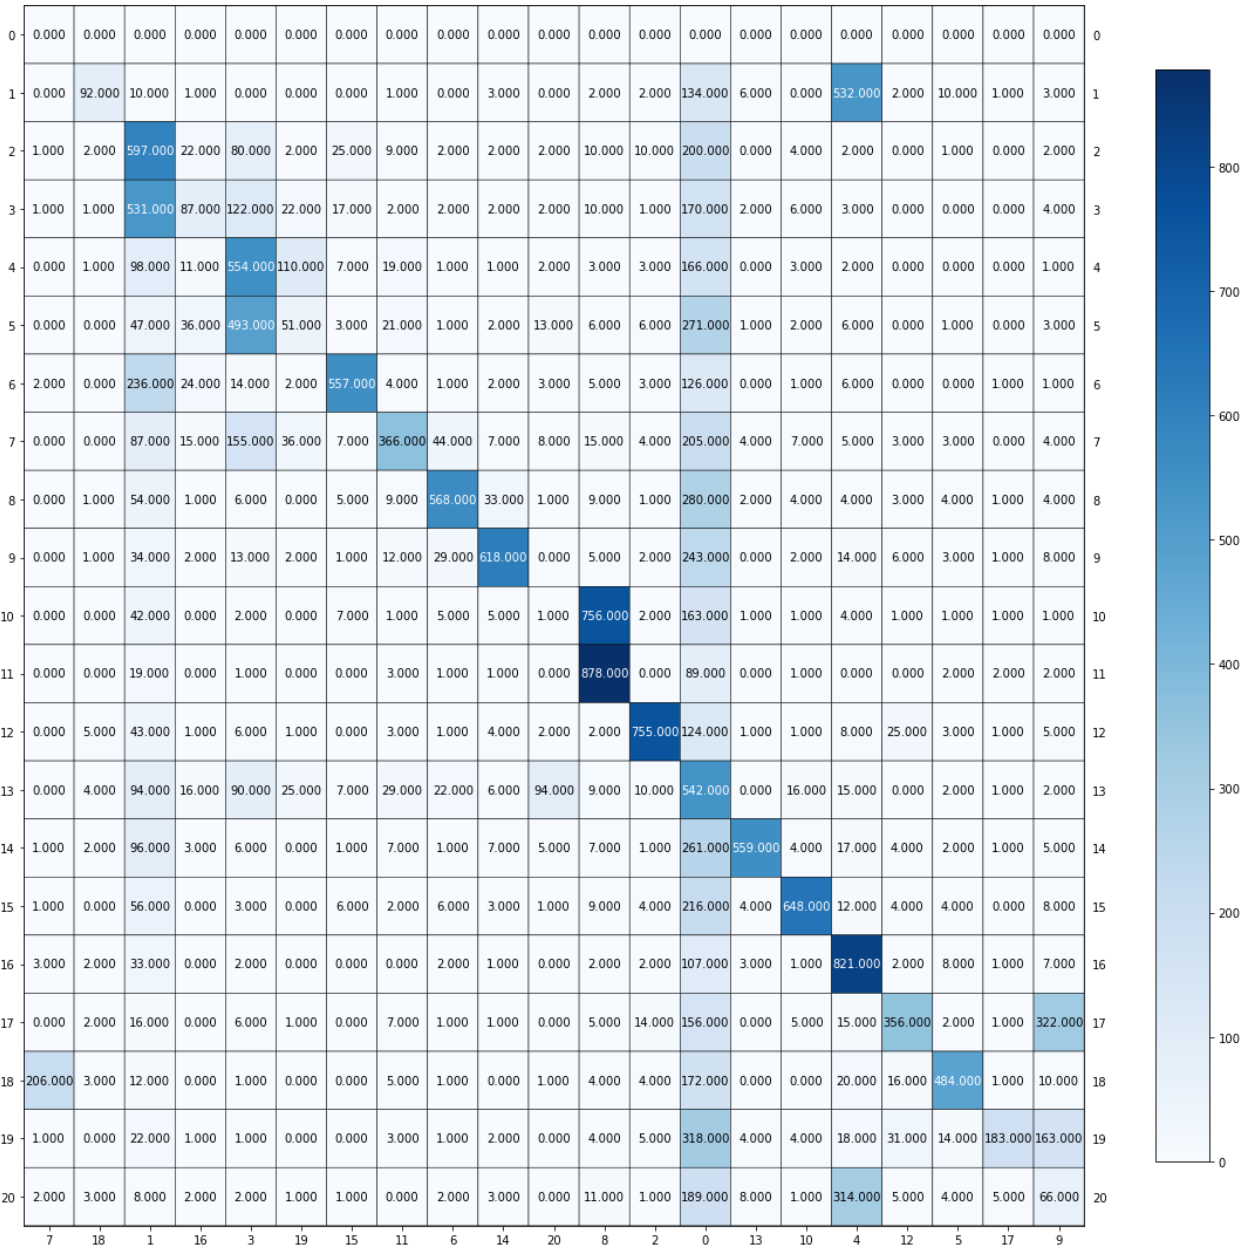
Completeness Score 0.5282599394438849

V-measure Score 0.4955310283931022

Adjusted Rand Index 0.2483640672034359

Adjusted Mutual Information Score 0.4936991642480833

metric is: euclidean



For HDBSCAN we tested metrics, cluster_selection_epsilon and min_samples

eps: 0.2

min_sample: 150

Homogeneity Score 0.3750396494844003

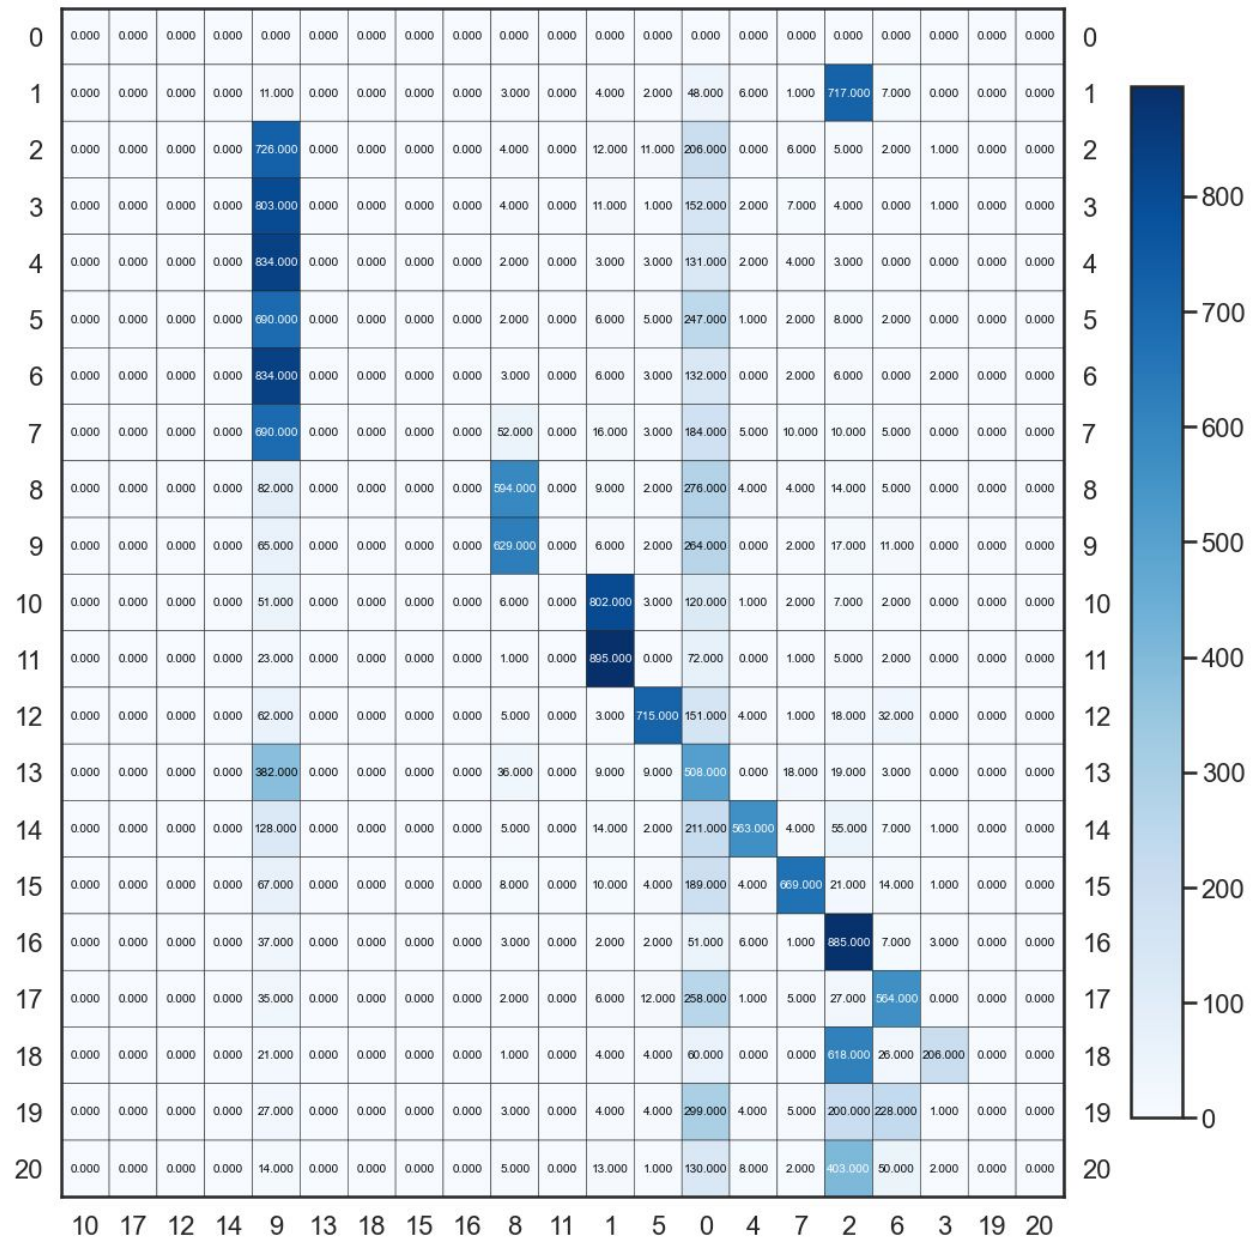
Completeness Score 0.5701629753929293

V-measure Score 0.45246112698448915

Adjusted Rand Index 0.19742645616795113

Adjusted Mutual Information Score 0.4514521062280271

metric is: canberra



Question 15:

There are 21 clusters in the DBSCAN model and 10 clusters in HDBSCAN model. HDBSCAN is a noise aware cluster algorithm. The -1 label are the data samples that are not assigned to any clusters, -1 results are outliers,

Question 16:

1. For this part, I downloaded the datasets on my PC and used the `pandas.read_csv` function to store the data in a pandas dataframe.

2. Then I changed the 'Category' column into categorical objects and created a new column named 'Category ID' which corresponds to the 5 categories. 0 is business, 1 is entertainment, 2 is politics, 3 is sport and 4 is tech.
3. The next step is to transform the Text column into TF-IDF vectors.
4. First, I tried to use SVD to reduce data and use KMeans() clustering. I tested several r from 1 to 300. However, the best V-measure score I got is only around 0.235.
5. And I decided to use the UMAP with parameters n_components = 100, metric = 'cosine' to reduce the dimensionality of the TF-IDF matrix. Then use Kmeans clustering method with parameters KMeans(n_clusters=5, random_state=0, max_iter=1000, n_init=30), we reached a very decent result.

Homogeneity Score 0.7708326965955365

Completeness Score 0.7715061495302332

V-measure Score 0.7711692760333082

Adjusted Rand Index 0.8046862093949397

Adjusted Mutual Information Score 0.7703963769832721

